MCS: A Formal Framework for Mathematical Knowledge and Learning Routes with External Learner Models

Peiheng Liu
peihengmath@gmail.com

October 1, 2026

Abstract

Relations in mathematical learning resources may encode joint requirements, alternative entries, proof dependencies, or presentation choices. These meanings must remain distinguishable when a learner model selects a route or a view hides part of its context. We present MCS, a specification separating versioned public mathematical resources, external learner models, and derived views and event routes. Resource ports retain evidence kind, theory, assumptions, version, and provenance. We prove a conditional saturation result for concept occurrence aggregated over unrestricted equivalent expressions; distinguish route-relative requirements from alternative sufficient supports; and give a compositional resource-preservation theorem and a bounded exhaustive-search guarantee. A lossless encoding into an annotated incidence graph makes the representational boundary explicit. A standard-library Python artifact replays local mathematical certificates, finite comparison experiments, and a limit-based route example. On the tested fragment, typed action graphs, annotated bipartite graphs, and MCS give identical classifications. The contribution is a precise, auditable integration of interfaces and their guarantees; the results establish neither unique graph expressiveness nor empirical learning benefits. General correctness of the Python proof checker and complete machine certification of the mathematical case studies remain open.

1 Introduction

A mathematical explanation rarely starts and ends with a node and an arrow. A proof can require two results together; a definition can be approached through two representations; and a learner can revisit the same concept in two different events. Moreover, reading a theorem, obtaining its proof, and being able to apply it are distinct outcomes. A planning system should say which one it has produced.

Two motivations lead to the present specification. Within mathematics, an inference needs the appropriate theory, open assumptions, and evidence, even when intermediate material is hidden in a compact view. In learning support, different people need different entry points and may leave some prerequisites unconfirmed. Their personal records should not silently alter the shared mathematical content. The first motivation concerns the legitimacy of a resource connection; the second concerns the ownership and interpretation of a decision.

Consider two actions producing a goal g:

$$a_1 : \{u, v\} \longrightarrow \{g\}, \qquad a_2 : \{w\} \longrightarrow \{g\}.$$

If every incoming node edge is treated as required, all of u, v, w are demanded. If every edge is treated as sufficient, u alone appears enough. Explicit action vertices and grouped input ports resolve this ambiguity. They do not by themselves resolve another one: an output describing g is not necessarily a proof of g. Types, assumptions, versions, and sources must also be matched. Annotated graph models can retain this information.

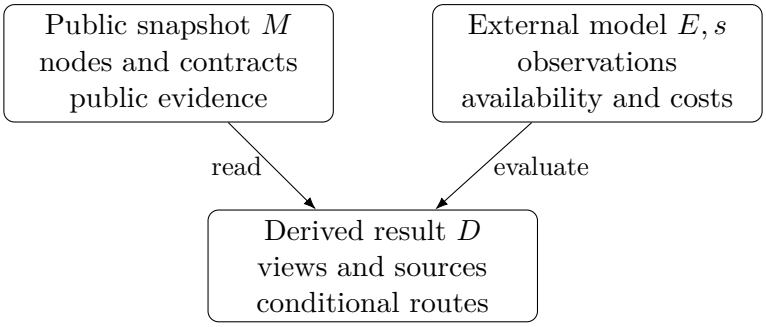


Figure 1: A dependency and ownership contract. Public revision is a separate, explicit operation. The diagram is a specification, not evidence of an implementation's isolation or a learner model's validity.

The construction develops from entries and links, through grouped actions and explicit event sources, to a separation of public resources and personal evaluation. MCS denotes this integrated specification, rather than a new logic or a psychological law. It supports implementations whose output states both what a route supplies and why the route was accepted.

Contributions and their status. The paper makes three connected contributions.

1. It specifies a resource and event interface that keeps mathematical evidence, displayed content, and external learner availability separate. The principal guarantee is compositional resource preservation under explicit local validity and source-matching hypotheses (Theorem 6.2).
2. It collects the mathematical boundaries needed by that interface: unrestricted equivalence can destroy the discriminatory value of concept occurrence (Theorem 5.2); common requirements do not represent alternative sufficiency; and exhaustive planning is relatively complete only over an explicitly enumerable finite domain (Theorem 8.2).
3. It supplies a portable finite artifact, a running mathematical example, and an equal-information comparison. A constructive incidence-graph encoding and the finite results identify where no representation advantage has been established (Proposition 7.2).

Several ingredients are elementary applications of established mathematics. Topological sorting, intersection antitonicity, read-only invariance, and exhaustive enumeration are not presented as new general theorems. The proposed contribution lies in their explicit integration with evidence and scope obligations. No claim of priority for the individual counterexamples is made.

Evidence vocabulary. DEF denotes a definition or modeling choice; PROOF an argument given here; REF an attributed external result or design; FINITE an executed check on a stated finite domain; ILLUSTRATION stipulated example data; and NOT-CLAIMED a result not established here. These labels answer different questions and are not a ranking of confidence. A passed program run is not a proof that the program is correct on every input.

2 Related work and intended scope

Mathematical knowledge representation. OMDoc represents mathematical material at document, statement, and theory levels (OMDoc Project, n.d.). MMT provides a modular representation of theories and theory morphisms while separating foundation-dependent and foundation-independent concerns (Rabe and Kohlhase, 2011). These are direct points of comparison. MCS resource payloads and evidence can in principle refer to such libraries. This

paper specifies the interface from public resources to condition-bearing learning events; it does not implement an OMDoc/MMT translation or prove a stronger representation theorem than those systems. An MMT *view* is a theory morphism, whereas a view in Section 7 is a display projection with a boundary.

Knowledge and learning spaces. Doignon and Falmagne (2015) distinguish knowledge spaces from the quasi-ordinal spaces that are additionally closed under intersection; their Definition 9 and Theorem 10 relate the latter to quasi-orders. We use this existing distinction to situate the small alternative-entry example in Section 5. MCS does not assume that every learner follows a union-closed state family, that learning is irreversible, or that mathematical nodes coincide with assessment items. Such assumptions belong to a particular external model.

Partial-order planning. The POCL formulation described by Younes and Simmons (2003) includes action instances, causal links supplying conditions, and ordering constraints; it permits multiple instances of one action schema. Our event identities and source references have clear counterparts there. Public resources in the present core are persistent and reusable: there are no delete effects, so the usual threat problem for consumed or negated conditions is outside this core. Forgetting and limited capacity belong to the external schedule predicate. The exhaustive procedure below is a reference specification, not a competitive POCL planner.

Adaptive educational systems. Brusilovsky (1997) describes systems combining domain models, student models, learning material, and adaptive teaching operations. Separating public content from student information therefore is not in itself a new architecture. Here the focus is the status of resources passed between actions, the scope of a prerequisite assertion, and the conditions under which a display transformation preserves a route certificate. We do not infer educational effectiveness from the existence of these interfaces.

The comparison question. The relevant question is whether an implementation detects an invalid evidence connection, preserves unresolved conditions, and reconstructs its decision. Calling a data structure a graph or a space does not settle that question. We compare encodings with equivalent annotations and retain negative and equal results. Formal interoperability with the above systems, maintenance benefits at scale, and classroom outcomes require further work.

3 A running example: two approaches to a limit

Let $D \subseteq \mathbb{R}$, let a be an accumulation point of D, and let $f : D \to \mathbb{R}$. There are two familiar ways to state $f(x) \to L$ as $x \to a$: a punctured-neighborhood condition and a condition on every approaching sequence. A learner comfortable with inequalities may first inspect the former; a learner familiar with sequences may first inspect the latter. These are presentation choices, not two different truths.

Definition 3.1 (The two statements; DEF)**.** *Write* $\mathrm{ED}(f, a, L)$ *for*

$$\forall \varepsilon > 0 \ \exists \delta > 0 \ \forall x \in D : 0 < |x - a| < \delta \implies |f(x) - L| < \varepsilon.$$

Write $\mathrm{SEQ}(f, a, L)$ *for the assertion that every sequence* $(x_n)_{n \geq 1}$ *in* $D \setminus \{a\}$ *with* $x_n \to a$ *satisfies* $f(x_n) \to L$.

The puncture excludes the value at the center; the universal quantifier over sequences prevents one favorable sequence from certifying the limit. The accumulation-point condition ensures the usual nonvacuous setting and uniqueness. It is not needed solely for the equivalence below: at an isolated point both punctured formulations are vacuous.

Action	Joint inputs	Output and evidence
`define_ed`	Metric, Quantifiers	LimitED content; definition presentation
`define_seq`	Metric, Sequences	LimitSeq content; definition presentation
`bridge`	LimitED, LimitSeq	Bridge reviewed-proof; Proposition 3.2
`show_bridge`	None	Bridge content; statement presentation only

Table 1: The executable presentation/reviewed-proof fragment. All ports retain an explicit real-analysis context. The bridge action records the written proof as a trust input; it does not mechanically derive it.

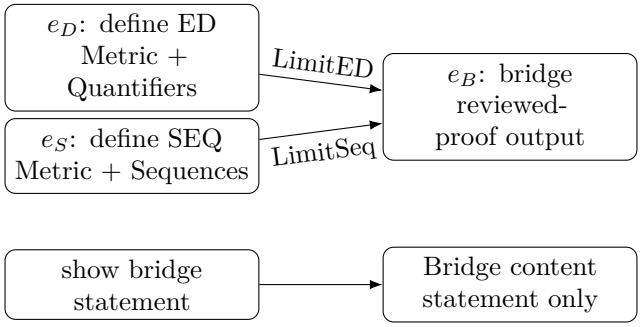


Figure 2: The running case has joint inputs to the bridge event, with either definition presented first. Showing the same statement separately produces content, which cannot fill the reviewed-proof target.

Proposition 3.2 (The limit bridge; PROOF)**.** *In the stated real metric setting and ZFC metatheory,* $\mathrm{ED}(f, a, L)$ *holds if and only if* $\mathrm{SEQ}(f, a, L)$ *holds.*

Proof. The forward direction applies one neighborhood bound to a sequence tail. The reverse constructs a bad sequence if neighborhood control fails. Assume ED, and fix an admissible sequence and $\varepsilon > 0$. Choose $\delta > 0$ as in the definition. Eventually $|x_n - a| < \delta$; since $x_n \neq a$, eventually $|f(x_n) - L| < \varepsilon$. Conversely, failure of ED gives $\varepsilon_0 > 0$ such that for every $\delta > 0$ there is an $x \in D$ with $0 < |x - a| < \delta$ and $|f(x) - L| \geq \varepsilon_0$. For $n \geq 1$, choose such an x_n with $\delta = 1/n$. This countable choice is available in the stated metatheory. Then $x_n \to a$, but $f(x_n) \not\to L$, a contradiction. □

For continuity assume additionally $a \in D$. Continuity at a implies the punctured condition with $L = f(a)$. Conversely that condition handles $x \neq a$, and the missing case $x = a$ has zero error. This proves the continuity bridge in this setting.

It is easy to erase a condition while preserving a plausible picture. Checking just positive sequences at zero would accept a function equal to 1 on positive reals and -1 on negative reals as having limit 1. Allowing sequences that repeatedly equal the center tests its value too: if $f(0) = 1$ and $f(x) = 0$ for $x \neq 0$, the punctured limit is 0, but the constant zero sequence does not have image limit 0.

The public library records both definitions and the written bridge proof. A presentation of the bridge statement has a different evidence kind from the reviewed proof, and both differ from a machine-checked proof object. The reference route example has four registered actions (Table 1). The label reviewed records the level of evidence; it does not assert that external peer review has occurred.

4 Public resources and external evaluation

4.1 A small public interface

We work in ordinary set theory. Records and certificates used by an implementation have finite descriptions; the general public library need not be finite. The route results require a resource-checking interface rather than a particular complete logical calculus.

Definition 4.1 (Resource descriptor; DEF)**.** *A resource descriptor is a tuple*

$$r = (n, v, \theta, \kappa, \Delta),$$

where n is a node identifier, v its payload version, θ a theory or context identifier, κ an evidence kind, and Δ a declared finite set of open conditions. A source occurrence is a descriptor together with a provenance reference. The strict interface matches descriptors by equality.

Examples of κ are content, declared-assumption, reviewed-proof, and checked-proof. Equality of node identifiers alone is insufficient. Conditions are not discharged merely because a source is available. A checked proof under Δ remains conditional on Δ. A learner's belief is an external observation, not a proof-resource kind.

Strict matching rejects some legitimate transports, including weakening to a larger assumption context or moving through a theory translation. Such transports can be explicit actions with their own evidence. We do not silently make them part of node equality.

Definition 4.2 (Public snapshot; DEF)**.** *A public snapshot M contains immutable versioned node payloads, their representations, a resource registry, public evidence records, and a family Λ_M of action contracts. Identifiers resolve within a declared version boundary. Public revision produces a new snapshot or an explicit version map and is separate from personal-state updates.*

This is the part of the larger MCS organization needed here. Topics, complete template catalogs, and specialized views can be layered on top. They are not additional hypotheses of the route theorems.

Definition 4.3 (Action contract; DEF)**.** *A contract $a \in \Lambda_M$ consists of a finite indexed input family I_a, a finite indexed output family O_a, a mode, and local evidence w_a. Indexed ports allow equal descriptors to have distinct declared roles. The local validity judgment* $\mathrm{Valid}_M(a, w_a)$ *is specified by the mode and is independent of the existence of a route.*

For a proof-producing action, local evidence accounts for the stated theory, conclusion, and open conditions. For a presentation action it accounts for supplied content and its reference. In the finite case study a written proof is an explicit trusted input. Only the separately replayed certificate fragment is labeled checked-proof.

Logical boundary. A checker may accept a certificate p for $T; \Gamma \vdash \varphi$. Acceptance establishes that judgment only relative to correctness of its checking rules and their implementation. Truth in a model additionally requires soundness and satisfaction of T, Γ. The main route theorem is parameterized by local validity and does not prove soundness of a Python HOL implementation. It makes no completeness claim for standard higher-order semantics.

4.2 What is external

Definition 4.4 (External model and result; DEF)**.** *An external model E declares a state space X_E, a binding to the public version, observation and transition interfaces, and a schedule-evaluation predicate or procedure. A state $s \in X_E$, goal g, budget b, and policy ρ are explicit inputs. A derived result*

$$D(M, E, s, g, b, \rho)$$

contains views, candidate event routes, evaluation results, and provenance. It does not implicitly update M.

General misconception patterns and assessment items may be public. Who exhibited a misconception or answered an item belongs to E. A goal separates its public resource target from a capability claim evaluated by E. One correct answer need not establish every capability attached to a node.

Unknown information is explicit:

$$\mathrm{Info}(A) = \{\mathsf{Unknown}\} \sqcup \{\mathsf{Known}(a) : a \in A\}.$$

For supports, $\mathsf{Known}(\varnothing)$ means a precisely known empty set; $\mathsf{Unknown}$ means the support has not been established. For learner input, unknown availability is neither confirmed availability nor confirmed unavailability. Missing answers are not stored as wrong answers.

Proposition 4.5 (Read-only invariance; PROOF)**.** *If every operation used to compute D reads a fixed snapshot M and writes only external or derived records, changing E, s, g, b, ρ leaves the snapshot and its registered public relations unchanged.*

Proof. Each operation preserves the public record by the write restriction; induction on any finite execution gives the result. □

The premise is an implementation obligation. Placing personal data in another table does not establish it. Nor does the proposition compare truth in two different mathematical models of a theory.

5 Concept occurrence, support, and alternative requirements

5.1 A failure of unrestricted equivalence aggregation

An expression can mention an idea without depending on it. Fix a simply typed language with lambda abstraction and application. A registry K assigns each concept c a closed typed representation t_c. Let κ map alpha-normalized closed-term codes to finite sets of registered concepts, with $c \in \kappa(t_c)$. This is a naming convention, not a semantic concept-discovery algorithm.

Definition 5.1 (Occurrence and equivalence aggregation; DEF)**.** *For a finite typed expression e, inspect its unreduced syntax tree. At each position close the subterm over its free variables in a fixed context order and apply κ. Let $C_{\mathrm{occ}}(e)$ be the union of these finite registry answers. For initial representations $\mathrm{Rep}_0(n)$, define*

$$C_{\mathrm{all}}(n) = \bigcup_{e_0 \in \mathrm{Rep}_0(n)} \bigcup_{e \equiv e_0} C_{\mathrm{occ}}(e),$$

where equivalence keeps the same language, type, theory, and open context.

Theorem 5.2 (Conditional saturation; PROOF)**.** *Suppose $\mathrm{Rep}_0(n) \neq \varnothing$, equivalence contains ordinary beta-conversion, and occurrence scans the unreduced syntax. Then $C_{\mathrm{all}}(n) = K$.*

Proof. Mention any concept in an unused argument while preserving the expression's value. Fix $e_0 : \tau$ and $c \in K$, with $t_c : \beta$. For fresh $z : \beta$,

$$e_c = (\lambda z^{\beta}.e_0)t_c \equiv e_0.$$

The closed subterm t_c occurs before reduction, so $c \in C_{\mathrm{occ}}(e_c)$. Thus $K \subseteq C_{\mathrm{all}}(n)$. The reverse inclusion follows from the registry's codomain. □

This concerns the specified function, not all semantic dependency analyses. With no initial representation the union is empty. If transformations are only identities, or occurrence is computed after beta-normalization, this argument fails: a closed constant d need not mention a distinct registered constant c. Unnamed semantic objects are outside the registry and theorem. The result does not recover all concepts from all mathematical meanings.

5.2 Supports record a use

Replace the unrestricted union by an indexed family. A support index specifies a representation or proof or route, its purpose, context, version, and witness. The support records what that particular object uses, not global necessity. A proof may contain a redundant premise, and two valid proofs may use different lemmas. For a fixed purpose u, let

$$\mathcal{A}_u(n) = \{A : \exists w \ \ \mathrm{Verify}_u(n, A, w)\},$$

$$\mathrm{MinSupp}_u(n) = \{A \in \mathcal{A}_u(n) : \nexists B \in \mathcal{A}_u(n),\ B \subsetneq A\}.$$

The witness relation must be independently defined. Helping someone understand is an external empirical predicate, not an unqualified instance of this public relation.

Proposition 5.3 (Minimal supports; PROOF)**.** *If $\mathcal{A}_u(n)$ contains a finite member, it contains an inclusion-minimal member. If the candidate domain is finite and adequacy is decidable, all minimal supports can be enumerated.*

Proof. Among adequate subsets of one finite adequate set choose one of minimum cardinality. An adequate proper subset would contradict that choice. For computation enumerate all candidate subsets and remove each one containing an adequate proper subset. □

Minimality need not be unique: the sufficient supports $\{a\}$ and $\{b\}$ give alternatives. Infinite adequacy families need not have minimal members; the cofinite subsets of $\mathbb{N}$ are an example. Inclusion-minimality and minimum cardinality are different objectives.

Safe unknown support. Interpret $\mathsf{Unknown}$ as all subsets of K, and $\mathsf{Known}(A)$ as the single possibility A. A set operation returns a precise result only if all possible inputs give that result. Hence $\mathsf{Unknown} \cap \mathsf{Known}(\varnothing) = \mathsf{Known}(\varnothing)$, but in general $\mathsf{Unknown} \cap \mathsf{Known}(\{c\}) = \mathsf{Unknown}$. Bounds $L \subseteq A \subseteq U$ can use a separate interval interface; a nonexact interval must not be labeled an exact support.

5.3 Relative requirements are not a sufficient starting set

After a route is defined, let $\mathsf{Req}_r(g)$ collect input descriptors of events active in producing g, omitting irrelevant branches. For a declared nonempty route family $\mathcal{R}$, define

$$\mathrm{Common}_{\mathcal{R}}(g) = \bigcap_{r \in \mathcal{R}} \mathsf{Req}_r(g).$$

This depends on the library, background, goal, and route family. For $\varnothing \neq \mathcal{R}_1 \subseteq \mathcal{R}_2$,

$$\mathrm{Common}_{\mathcal{R}_2}(g) \subseteq \mathrm{Common}_{\mathcal{R}_1}(g),$$

since membership on the left holds for every route in the larger family. A personal filter can increase common requirements; its result cannot be written back as a universal public prerequisite.

For supports $\{a\}$ and $\{b\}$, the common set is empty, but neither route runs without an entry. Their union incorrectly requires both. An empty family reports absence of a route in that scope, rather than using vacuous universal quantification to mark every resource necessary. Even a common input can be produced on the way: $\varnothing \to \{m\} \to \{g\}$ uses m without needing it initially. Initial requirements need a separate analysis of background sources.

Example 5.4 (A single item order loses an alternative). On $Q = \{a, b, c\}$, consider

$$\mathcal{K} = \{\varnothing, \{a\}, \{b\}, \{a, b\}, \{a, c\}, \{b, c\}, \{a, b, c\}\}.$$

It is union-closed and accessible by single-item additions, but not intersection-closed: $\{a, c\} \cap \{b, c\} = \{c\} \notin \mathcal{K}$. All downsets of one partial order are intersection-closed; thus $\mathcal{K}$ is not that family for an order on these three items. This instantiates the established distinction discussed by Doignon and Falmagne (2015), not a limitation of arbitrary annotated graphs.

6 Event routes and preservation of resource obligations

The same action may be used twice, so precedence belongs to events. A source must identify which occurrence supplies which input port. This is where an informal path becomes a checkable resource connection.

Definition 6.1 (A sourced event route; DEF)**.** *Fix M, a finite indexed background B, and a finite target family G. A route contains:*

1. *a finite event set P and a contract label* $\mathrm{act}(p)$*;*
2. *a source for every event input port, either a particular background port or a particular output port of another event, with equal descriptors;*
3. *a source of equal descriptor for every target port;*
4. *optional policy-order edges K_P.*

An event source q for an input of p creates an edge $q \to p$. The union of all source edges and K_P must be acyclic. Every used contract has accepted local evidence. Background sources retain their context and provenance. Resources persist and may be reused. The reflexive transitive closure of the combined graph is $\preceq$.

The result package retains all these data. Its poset isomorphism type alone is insufficient: two two-event chains can supply different resources. Renaming events must rename all source and policy endpoints consistently.

Theorem 6.2 (Compositional preservation; PROOF)**.** *Let r satisfy Definition 6.1. Then:*

1. *$(P, \preceq)$ is a finite partial order, and every linear extension provides every event input from its designated earlier producer or background. All target ports retain matching sources.*
2. *Fix an interpretation of resource judgments. If background judgments hold and every used contract locally preserves those judgments, then all event-output and target judgments hold after their producers execute.*
3. *These statements remain true after a display restriction retaining the same resolvable contracts, sources, and hidden-resource boundary.*

Proof. First order events, then propagate the exact resource judgments. Reflexivity and transitivity come from closure; mutual reachability of distinct events would give a cycle, so antisymmetry follows. A finite acyclic graph has a topological ordering, obtained by repeatedly removing a vertex with no incoming edge. In any linear extension, the chosen producer q of an input to p precedes p.

Induct on the position of p. Background sources are available initially. Event sources were produced at earlier steps and persist. Matching preserves the complete descriptor, including kind, context, conditions, and version. The local preservation premise therefore applies to the required input judgments and yields the output judgments. Target references use the same argument or a background source. Finally, a display restriction satisfying the retention condition changes no source in this induction, so the same proof applies. □

Part 1 is structural. Part 2 depends on the meaning and soundness of local validity. It cannot turn a displayed proof into a verified certificate, discharge an assumption, or conclude learner mastery. Part 3 is a conditional display guarantee, not a claim that arbitrary subgraphs preserve consequence.

Why the hypotheses matter.

- Matching only n can connect Bridge content to a required Bridge proof. Dropping Δ can turn a conditional claim into an unconditional one; dropping θ or v can accept a result from another theory or an obsolete payload.
- Source edges $p \to q$ and a policy edge $q \to p$ are separately acyclic but cyclic together. Checking only the source graph is insufficient.
- Without local validity, the contract $\varnothing \to \{\text{false theorem}\}$ defeats the semantic conclusion while remaining a well-formed graph.
- If resources are consumable, one token cannot supply two consuming events without a quantity discipline. The persistent-resource induction fails.
- Removing a hidden input instead of retaining its environment reference prevents that source from justifying execution.

Finiteness gives a finite schedule. Without it, an acyclic chain indexed by the negative integers, each requiring its predecessor, can lack a first event.

Structural order and personal order. Independent events may be swapped structurally, yet fatigue, forgetting, or explanation order may change an external evaluation. The result retains a chosen schedule and its evaluation. Structural commutativity is not psychological equivalence.

7 Views and an equal-information graph encoding

Definition 7.1 (View with boundary; DEF)**.** *A restriction view selects visible resources or nodes and retains provenance to M. Its boundary contains the referenced resources hidden by the selection, with sufficient versioned references to resolve them. Checks use M, not an assumption that the visible induced graph is the entire environment.*

The artifact computes the boundary for a given finite route by visiting input and target references. A view supporting every contract in a larger library needs a boundary for that declared scope. Aggregation additionally needs a map from displayed blocks to members; it is outside the executable restriction example.

For predicates P, Q on a fixed snapshot, pure restrictions select $P \cap Q = Q \cap P$ if the boundary is recomputed from final references. General operators need not commute. In $a \to b \to c$, dependency completion from c, followed by restriction to $\{a, c\}$, retains a, c and the missing dependency in the boundary. Restricting first and computing closure only in the induced graph loses the connection through b. Local visibility and self-containment are different properties.

Proposition 7.2 (Lossless incidence encoding; PROOF)**.** *A finite public-contract instance admits an annotated bipartite incidence encoding with an inverse decoding on well-formed encodings. Sourced routes and the acceptance conditions of Definition 6.1 are preserved.*

Proof. Create a resource vertex for each distinct descriptor, with all its fields, and an action vertex for each contract, with its mode and evidence. Use indexed input and output incidences

for every port, including repeated descriptors. Keep event identities separate from action vertices. Encode event action references, input sources, target sources, and policy edges without changing their indices.

Decoding reads these fields and orders ports by their index. On well-formed encodings it recovers the records up to consistent graph-vertex renaming. Descriptor equality, local validity, the combined event graph, and target provenance are unchanged. The same external predicate receives the same schedule and labels, so its decisions are unchanged. □

Port indices and evidence fields are essential to this inverse; an unannotated node graph need not retain them. This precludes an expressiveness advantage over the equal-information incidence representation. Differences in usability, maintenance, or solver performance require measurement.

8 Bounded planning and failure states

Planning chooses actions and sources before ordering events. Each action's inputs are conjunctive; alternative actions and sources are separate candidates. Combining all alternatives into one precedence graph changes the problem.

Definition 8.1 (Finite planning domain; DEF)**.** *Fix a finite contract list, finite background and target candidate lists, an event bound h, and finite, effectively listable domains for every annotation, policy-order choice, and external constraint instance. Contract validity, descriptor matching, and the external predicate $\chi(E, s, g, b, \rho, r, \ell)$ for route r and linear extension ℓ are decidable. Supplied certificates are fixed inputs; if searched for, their syntax and size have explicit finite bounds.*

This is the domain of the completeness assertion. A finite number of nodes does not bound repeated events, proof lengths, continuous parameters, or the complexity of external probability calculations.

Reference procedure. Enumerate action words of lengths $0, \dots, h$, with event names $1, \dots, k$. For every input enumerate matching background or earlier output sources. Enumerate every target source, annotation, and declared policy/constraint choice. Reject invalid contracts and cyclic combined graphs. For each remaining graph enumerate its linear extensions and apply χ. Retain every accepted schedule with full route data. Deduplicate only under isomorphisms preserving all externally relevant data. Pareto filtering or display selection comes later.

Theorem 8.2 (Bounded search; PROOF)**.** *Under Definition 8.1, the procedure terminates, accepts only candidates satisfying the declared conditions, and retains a representative of every feasible route and schedule in the bounded domain, up to consistent event renaming.*

Proof. For m actions there are $\sum_{k=0}^{h} m^k$ words. Every finite word has finitely many input and target ports, each with finitely many sources. The remaining lists are finite and completely enumerable. There are at most $k!$ linear extensions for k events; each check terminates. This proves termination.

Acceptance checks every declared condition, so accepted routes satisfy the resource theorem and chosen external predicate. Conversely, a feasible bounded route has a topological ordering of its combined graph. This ordering occurs as an enumerated action word after event renaming. Its sources, targets, annotations, policy constraints, and accepted schedules occur in the complete lists. A representative is retained. The stipulated deduplication cannot identify externally distinguishable candidates. □

A coarse bound. Let each action have at most d inputs and o outputs, the background b_0 ports, and the target q ports. Let A_k bound other annotation and policy choices for a word of length k, let $S_k = \max(1, b_0 + ko)$, and let C_k bound checking time for a fully specified schedule. Up to construction overhead, enumeration is bounded by

$$\sum_{k=0}^{h} m^k S_k^{kd+q} A_k \, k! \, C_k.$$

This overcounts incompatible sources and duplicate event renamings, and makes no tractability claim. For $m = 0$, the length-zero word remains the single empty candidate.

Statuses under incomplete information. The strict theorem uses decisive checks. A practical interface additionally distinguishes:

- Found: a route and schedule meet all declared, confirmed conditions within the selected model;
- Conditional: structural data are checked but an explicitly listed entry or model assumption is not confirmed;
- InfeasibleWithinBound: every candidate in the declared bound is decisively rejected and none is pending;
- Unknown: search, evidence resolution, or a required decision remains incomplete or unsupported.

A timeout cannot justify infeasibility. Conditional candidates cannot be discarded to produce a false infeasibility result. In the executable fragment, availability has three values and conditionality comes only from using an unknown entry. It does not realize every external predicate in the interface.

A computable real-valued probability alone need not allow a terminating exact threshold comparison. Certified bounds can decide when they separate the threshold; unresolved comparisons stay unknown. Optimality is relative to the model, objective, and bound. Unknown effects cannot be replaced by default numbers to manufacture Pareto dominance.

9 Implementation and finite evaluation

9.1 Artifact and trust boundary

The `anc/` directory contains Python standard-library code, frozen inputs, local certificates, and a one-command replay runner. Selected existing certification files are copied byte for byte and pinned by SHA-256. The runner works in a temporary copy and writes to a newly named directory; it does not refresh the evidence used as its own reference.

The mathematical checker and action adapter are Python implementations, not machine-verified kernels. Tests exercise specific cases and two controlled defective variants, not general soundness. The route module implements a smaller interface than Definition 8.1: persistent descriptors, bounded action words and source choices, fixed integer costs, and three-valued entry availability. Normal enumeration adds no policy edges or dynamic state transitions; the verifier separately tests policy edges and cycle rejection.

9.2 Local mathematics and the running route

The four checked fragments are: evaluation of a left-multiplication map at an identity; beta-reduction of a constant function; same-chart composition at a point under an inverse condition; and additivity of a rank-one map under stated algebraic assumptions. The limit certificate

Component	Executed scope	Outside the evidence
Certificate tests	28 checks, including 8 positive certificates and 19 negative certificate cases	General implementation soundness; full-calculus completeness
Independent semantics	35 terms / 175 comparisons; 42 case valuations; 11 background instances	All typed terms and all Henkin models
Action adapter	4 local mathematical fragments; 5 rejected source cases	Complete mathematical bridges
Representation study	2295 candidates per encoding	General representation or performance superiority
Limit route fragment	85 words at $h = 3$; 2 sourced routes; 11 negative cases	Learning success; a complete general planner

Table 2: Reproduced finite evidence. Counts refer to different domains and cannot be added into one coverage claim.

proves $(\lambda n.a)(n) = a$. It does not prove convergence, construct a real-number library, or prove Proposition 3.2. The artifact passes its accepted receipt into a one-event checked-proof route, demonstrating the certificate, resource, and route connection for this exact fragment.

In the running example, all three background content resources are available in the fully diagnosed profile. The target is reviewed Bridge. At $h = 3$ and unit costs, enumeration of $1 + 4 + 16 + 64 = 85$ words yields

$$(\texttt{define_ed}, \texttt{define_seq}, \texttt{bridge}) \quad \text{and} \quad (\texttt{define_seq}, \texttt{define_ed}, \texttt{bridge}).$$

Each has two linear extensions as a sourced partial order; enumeration retains both event-numbered records rather than quotienting them. An independent resource-accumulation check validates the four schedule instances. The `show_bridge` action cannot satisfy a reviewed-proof target.

Two stipulated profiles select different first presentations. Profile A confirms Metric and Quantifiers but leaves Sequences unknown; profile B confirms Metric and Sequences but leaves Quantifiers unknown. Both selected routes are conditional until the missing entry is resolved. Preferences are supplied parameters, not estimated advantages. A cold start retains conditional routes with all entries unknown. Hash checks confirm unchanged frozen public contracts in this run.

If only Bridge is visible, Metric, Quantifiers, Sequences, LimitED, and LimitSeq remain in the boundary. Dropping it is rejected. Other rejected cases include a changed version, theory, assumptions or evidence kind, a missing input port, a dangling goal, an inadequate budget, an unavailable entry, and a policy edge closing a source cycle.

9.3 Equal-information comparison

The inherited study fixes four nodes, four actions, horizon and budget three, and 27 three-valued initial-information states. The 85 words give 2295 candidates per encoding. A separate reference routine checks Boolean completions of unknown entries. Typed action, annotated bipartite, and MCS encodings decode to the same semantics and use the same evaluator. This checks encoding consistency, not the quality of independent planners.

The plain baseline fails because of its declared information loss; graphs in general need not fail. Enhanced encodings match exactly. The new example also round-trips all descriptor and evidence fields through an annotated incidence encoding and reproduces its route results. Neither experiment establishes a unique MCS advantage.

Encoding	Legal-route misses	Unjustified Found	Conditional errors
Plain node-dependency graph	140	2	130
Typed action graph	0	0	0
Annotated bipartite graph	0	0	0
MCS fragment	0	0	0

Table 3: Executed comparison. The plain baseline requires all incoming edges and omits grouping and resource types. Error columns are separate metrics, not necessarily disjoint sets.

Event bound	Action words	Sourced route records	Status
0	1	0	InfeasibleWithinBound
1	5	0	InfeasibleWithinBound
2	21	0	InfeasibleWithinBound
3	85	2	Found
4	341	30	Found

Table 4: Four-action limit fragment, with budget equal to event bound. Counts retain repeats, unused actions and distinct sources; they are not isomorphism-class counts or a speed benchmark.

9.4 A small enumeration growth check

This checks the stated bound in a small executable case, not asymptotic performance. The inherited synthetic learner-data exercise is retained for reproducibility; its scores support no pedagogical claim here.

10 Limitations and research questions

An interface does not identify a learning mechanism. Fix public presentation actions a, b with the same displayed target. Two external models have states s_0, s_1, initial state s_0, and $\mathrm{Obs}(s_i) = i$. In both, a moves any state to s_1. In E_1, b also moves any state to s_1; in E_2, it moves any state to s_0. Every history using only a has identical observations, yet the next b produces opposite outcomes. Both respect the public/external separation. The records and interface do not determine the untried effect. More observations of a cannot resolve this example; suitable observations of b or additional justified model restrictions are needed.

Formalization and implementation. The limit bridge has an ordinary mathematical proof here, with metatheoretic choice stated. Its implementation-level formalization is incomplete. A verified-kernel integration must establish that translations and adapters preserve the chosen judgments; certificates from another logic cannot be imported without that work. The resource theorem assumes local soundness; the Python artifact has no general machine proof of that premise.

Representation and efficiency. MCS can be implemented as an annotated graph. Future evaluations should hold annotations and evidence fixed and measure identification of stale dependencies, explanations of failed routes, or propagation of a version update. Equal results support a common specification or an engineering convenience, not superior expressiveness. Efficient search needs justified pruning or a solver with separately checked guarantees.

Empirical work. Preferred explanation order, analogy costs, and the diagnostic meaning of an answer are model choices until measured. Real comparisons need fixed targets, comparable

content and time, justified measurements, and a design separating selection effects from intervention effects. Educational experiments are not premises of the structural theorems, but are necessary for claims about gains.

Common misreadings. A shown node does not discharge a proof obligation; an empty common requirement set does not supply an alternative entry; a topological ordering does not certify learning; and a finite report does not certify its checker. Each mistake moves a statement between levels without its conditions.

11 Conclusion

MCS organizes public mathematical resources, external learner evaluation, and derived routes through explicit resource and source obligations. The guarantees are conditional: valid local connections persist along an acyclic reusable-resource schedule, display restrictions preserve them when boundaries are retained, and exhaustive search covers a stated finite decidable domain.

The guiding method is to state what a transformation preserves and which evidence supports it. Padding exposes the limits of occurrence aggregation; source tracing supports schedule induction; equal-information encoding constrains representation claims; and same-history/different-outcome models expose cognitive inference limits. A useful next question is which observable maintenance or learning task benefits once an enhanced graph receives the same information.

Availability and tool use. The source includes the finite artifact in `anc/`; Appendix A gives its command and boundaries. Generative AI assistance (OpenAI Codex) was used in restructuring and drafting this manuscript and in preparing code and finite checks. This disclosure does not constitute independent verification of mathematics, references, or implementation. Responsibility for the submitted content remains with the named author.

A Reproduction and claim-to-evidence map

From the extracted source directory run with Python 3.10 or later:

```
python anc/run.py -output replay
```

The output directory must not already exist. The reference run used Python 3.14.7. Version 3.10 is a syntax/dependency target, not a claim that every intermediate release was tested. Only the standard library is needed; no personal learner data or network service is required.

The runner verifies upstream hashes, copies certification code to a temporary directory, executes three suites, then runs the route fragment. Stable reports exclude diagnostic wall-clock timings. A package check reruns the extracted source and compares stable results.

Claim	Evidence and entry	Boundary
Saturation	Theorem 5.2, PROOF	Fixed registry, unreduced syntax, beta-conversion.
Source preservation	Theorem 6.2, PROOF	Local validity, exact matching, persistence, resolved sources, combined acyclicity.
Bounded completeness	Theorem 8.2, PROOF	Every candidate dimension effectively finite; every check decidable. No general implementation verification.

Claim	Evidence and entry	Boundary
Graph correspondence	Proposition 7.2, PROOF; `route-report.json`, FINITE	Lossless fields and ports. No unique expressiveness claim.
Local certificates	`kernel-report.json`, `action-report.json`, FINITE	Four mathematical fragments, not the complete limit bridge.
Encoding experiment	`research-report.json`, FINITE	Four actions, 27 information states, horizon and budget three.
Running case	`route-report.json`, FINITE/ILLUSTRATION	Stipulated profiles and costs. Checked constant-value resource is separately linked to an accepted receipt.
Teaching benefit	NOT-CLAIMED	No real learners, intervention, or measured cognitive effect.

B Assumption audit

The audit records the roles of hypotheses, not their independent necessity in every alternative formulation. Some can be traded for other assumptions.

Assumption	Failure or boundary if removed	Role
Nonempty representations	Empty union despite a nonempty registry	Choosing e_0 in saturation.
Admitted padding	Identity-only transforms expose no new concept	Constructing e_c.
Pre-reduction occurrence	Unused arguments disappear before scanning	Recovering t_c from syntax.
Finite registry answers	Finite traversal alone need not compute a finite support	Effective occurrence queries, not the set-theoretic equality itself.
Nonempty route family	Universal necessity becomes vacuous	Meaning of common requirements.
Complete resource match	Versions, conditions or kinds are confused	Applying the right contract.
Local preservation	Zero-input action asserts a false theorem	Semantic preservation.
Resolvable background	A source name alone supplies no resource	Induction base.
Combined acyclicity	Opposite source and policy edges	Existence of a schedule.
Persistence	Two consumers overuse a single token	Availability of earlier outputs.
Finite event bound	A single action has unbounded repetitions	Finite enumeration.
Complete finite annotation lists	Unbounded real annotations or omitted choices	Termination and coverage.
Decidable external checks	A predicate encoding arbitrary halting need not terminate	Decisive checking.
All source assignments	Choosing one producer omits an externally distinct route	Relative completeness.

Assumption	Failure or boundary if removed	Role
Data-preserving deduplication	Equal poset shapes hide different actions	Preserving distinguishable candidates.
Explicit unknown	Unknown entry treated as true	Conditional output and cold start.

References

Peter Brusilovsky. Integrating hypermedia and intelligent tutoring technologies: From systems to authoring tools. In P. Kommers, A. Dovgiallo, V. Petrushin, and P. Brusilovsky, editors, *New Media and Telematic Technologies for Education in Eastern European Countries*, pages 129–140. Twente University Press, Enschede, 1997. URL `https://sites.pitt.edu/~peterb/papers/EWED.html`.

Jean-Paul Doignon and Jean-Claude Falmagne. Knowledge spaces and learning spaces. arXiv:1511.06757, 2015. URL `https://arxiv.org/abs/1511.06757`. Especially section 3, Definition 9 and Theorem 10.

OMDoc Project. The OMDoc format. Project documentation, n.d. URL `https://www.omdoc.org/format/`. Accessed October 1, 2026.

Florian Rabe and Michael Kohlhase. A scalable module system. arXiv:1105.0548, 2011. URL `https://arxiv.org/abs/1105.0548`. The arXiv manuscript; especially sections 3–5.

Håkan L. S. Younes and Reid G. Simmons. VHPOP: Versatile heuristic partial order planner. *Journal of Artificial Intelligence Research*, 20:405–430, 2003. URL `https://www.cs.cmu.edu/afs/cs/project/jair/pub/volume20/younes03a-html/index.html`.